

Experiment 7

Build Grad-CAM yourself, and use it to see where a model is looking.

You need these two things:

1. A trained model. Use any model that is already trained on ImageNet — for example VGG-16 or ResNet-50. You will not train anything today. The model is already done.

2. Four pictures. Pick them carefully:

- Picture 1: one clear object, like a dog or a bird
- Picture 2: one clear object again, but something different
- Picture 3: two or more of the same thing (for example, two dogs in one photo)
- Picture 4: a small object with a busy or messy background

You must write Grad-CAM yourself. Do not use a ready-made Grad-CAM function.